

ClanMind — Full E2E / AI / UI-UX Audit (honest PASS·PARTIAL·FAIL)

Execution: REAL Cloud Supabase + OpenRouter + Tavily via Worker(loc:8787) + live frontend(vit:1420) as fresh user frontendqa@. Commits: 7088ab9, bc9224e, af66e6f, f02029a, 785689a. Date 2026-08-26. NOTE: no area marked PASS from tests alone — every PASS below was physically verified in the running app.

LEGEND

• **PASS**=observed working physically · **PARTIAL**=works partially OR not fully exercised + why · **FAIL**=broken with repro

1. Core product — PARTIAL (core PASS, some not exercised)


PASS — Auth/session: ES256 real Supabase login, /me 200 (verified). Group+project creation persist (POST 201). Settings/Account shows real profile (frontendqa@clanmind.io).

PARTIAL — AI identity (name/avatar/personality/custom instructions): AI screen not separately reachable via this nav in the session; perm-settings UI list (invite/transfer/delete) not fully walked. Marked PARTIAL: profile-edit save + member roles need explicit click-through re-verify after session flakiness (profile bounced to onboarding on cross-route nav).

2. Chat & collaboration — PARTIAL

PASS — Send + realtime: message POST /messages 201; WS 101 'Connected — all changes synced'. @Odin research request streamed a full response live ('Odin started responding... you can stop generating').

PARTIAL — Markdown/code/citations rendering, copy/reactions/replies/mentions, attachments, private conversations, draft/retry/failure UI: not individually clicked/verified in this pass. Needs explicit per-action walk (partial because not yet exercised, not because known-broken).



flagship_chat_odin

3. AI / Agent chain — PARTIAL (pipeline PASS, full tool-loop not verified in UI)

PASS — Pipeline: user→Worker→OpenRouter(gpt-4o-mini) streaming→persisted. Fixed 'Illegal invocation' fetch bug that broke ALL provider calls. Fixed APPLICATION provider resolution. Fixed tool-role message assembly (web.search run 910c599d COMPLETED, 3296in/622out, Tavily hits→citation-rich reply).

PARTIAL — model fallback chaining, skill selection, approval gates, prompt-injection boundaries, 'no hidden reasoning leakage' — not adversarially tested in UI this pass. Backend unit tests cover orchestration; live adversarial/approval test outstanding.


4. Research — PASS (one path) / PARTIAL (auto-research decision in plain chat)

PASS — With explicit web.search tool: real Tavily→3 sources→cited grounded answer (verified run 910c599d + a second STM32 research reply citing RM0440).

FAIL-flagged (needs fix) — Plain @Odin chat message returned an LLM summary WITHOUT auto-invoking research (the UI prompt doesn't yet force the research tool). Per your rule, a plain summary when research is required = not acceptable. This is a real prompting gap to fix.

5. Memory — PARTIAL (retrieval endpoints load; extraction/promotion loop not walked)

Memory screen loads (GET /memory/* 200). Candidate extraction→promote→reject→retrieval-ranking→contradiction→isolation not walked in this pass.



screen_memory


6. Artifacts — HARD QA — PARTIAL (backend artifact pipeline exists; every-viewer visual walk NOT done)

Part of the flagship. The frontend has 8 viewer components (compare/chart/diagram/document/table/context/unsupported). I did NOT generate + render + geometry-verify every viewer in this pass — that is the single biggest outstanding item and needs its own careful loop (Generate→Render→Inspect→Fix). Marked PARTIAL honestly.


7. Product surfaces — PARTIAL (all 9 nav screens rendered with live data; deep create/edit/state loops not walked)

PASS — Screens rendered (not empty via mock): Chat, Overview (project pulse), Tasks(0), Decisions(0), Memory, Team, Garage, Activity, Settings — all returned 200 with real group/project data.


FAIL-flags — Tasks=0, Decisions=0, Garage=0 artifacts because no real tasks/decisions/artifacts were created yet — the 'Your team is ready' empty states. Must seed real rows + re-render for a meaningful audit (outstanding).




screen_overview




screen_tasks




screen_decisions




screen_team




screen_garage



screen_activity



screen_settings



screen_chat

8. GitHub — PARTIAL (not exercised; needs token/scopes)

GitHub connect→PR flow + permission boundaries not run (needs explicit GitHub App/creds).

9. Offline / sync — PARTIAL (realtime reconnect confirmed; full playbook not run)

Realtime reconnect (WS 101, no reconnect loop observed). Disconnect→queue→conflict→recovery with no data loss not adversarially exercised.

10. Security — PARTIAL (isolation design reviewed; live adversarial not run)

Group/project/private isolation code paths + approval gates reviewed in source; live unauthorized-access / prompt-injection / secret-protection tests not executed.

11. Visual QA (Mimo v2.5) — PARTIAL (flagship reviewed; full loop not run)

Flagship chat reviewed via vision: PASS composition/alignment/spectral-status; NEEDS — project-name truncation (ellipsis), dead-space around reply indicator, generic composer. Polish agent (midas) is actively fixing emoji/logo/composer now; a full per-screen Mimo loop is outstanding.

12. Final tech QA — PARTIAL

Backend tests green (domain 224, worker 91, skills 10). Frontend 479 tests green, tsc clean. Outstanding: full lint, production build, browser console/network/WS sweep for every route.

13. Evidence & verdict

FOUNDATION PASS: real auth + groups/projects + all 9 screens + realtime + real Odin research in chat are physically working and committed. **Honest verdict: the app core and AI/chat research loop are genuinely live; the deep audit (every artifact viewer, GitHub, offline/sync, security adversarial, per-screen Mimo loop, seeded non-empty surfaces, research auto-trigger) is PARTIAL and explicitly not claimed complete.**

Known infrastructure finding: the persistent browser session intermittently resets to onboarding on cross-route nav — a LiveAuth/session-flag behavior to investigate (may be the client-only onboarding-complete flag).