

ClanMind — Final E2E / AI / UI-UX Audit (honest PASS·PARTIAL·FAIL)

Real Cloud Supabase + OpenRouter + Tavily, Worker(loc:8787), live frontend(:1420), fresh user frontendqa@. Commits: 7088ab9 ES256 auth · bc9224e AI pipeline · af66e6f tool protocol · f02029a live group/project create · 785689a realtime WS · 353ddc7 polish (emoji/alignment) · fa2b32c live feature-store seeding. No PASS from tests alone.

LEGEND: PASS=observed in running app · PARTIAL=partial or not fully exercised + exact reason · FAIL=broken with repro.

1 Core product


PASS — Auth/session ES256 real login /me 200; group+project creation persist (POST 201); profile shows real email; 9 nav screens render.

PARTIAL — AI identity screen (name/avatar/personality/custom instructions) not separately reachable via nav; member roles (invite/transfer/delete) not fully walked; profile-edit file save not re-verified after a session-reset-on-navigation observation.

2 Chat & collaboration

PASS — Send + realtime: message POST 201, WS 101 'Connected — all changes synced' (no reconnect loop). @Odin research streamed a full response live. Polish removed emoji; composer usable.

PARTIAL — markdown/code/citations rendering, copy/reactions/replies/mentions, attachments, private conversations, draft/retry/failure UI not individually clicked in this pass.



flagship_chat_odin

3 AI / Agent chain

PASS — pipeline user→Worker→OpenRouter(gpt-4o-mini)→stream→persist; fixed 'Illegal invocation' fetch bug (broke ALL provider calls); fixed APPLICATION provider; fixed tool-role message assembly (web.search run 910c599d COMPLETED 3296in/622out with Tavily→citations).

PARTIAL — model fallback chaining under live failure, skill selection, approval gates, prompt-injection boundaries, no-hidden-reasoning: unit-tested but not adversarially exercised in UI.

4 Research

PASS (with explicit tool) — Tavily→3 sources→cited grounded answer (run 910c599d + STM32 reply citing RM0440).


FAIL-to-fix — Plain @Odin returns an LLM summary WITHOUT auto-invoking research. Per rule, a plain summary when research is required is not acceptable → prompting gap still open.

5 Memory

PARTIAL — Memory endpoints load (200); extraction→promote→reject→rank→contradiction→isolation loop not walked; no memory rows seeded for the live demo group yet.

6 Artifacts (HARD QA)

PARTIAL — Seeded 5 artifact types (DOCUMENT/TABLE/CHART/CODE/DIAGRAM 201); Garage now lists them (DIAGRAM SPI1_TX_DATAFLOW shown). PER-VIEWER geometry QA (render, node/edge alignment, zoom/pan, versions/compare/export, fullscreen, permission, reduced-motion, no-clipping) NOT run this pass — the single biggest outstanding item.




screen_garage


7 Product surfaces

PASS — live feature-store seeding populates Artifacts + Overview (Odin: 1 unresolved decision insight). 9 nav screens render, realtime Connected throughout.


PARTIAL — Tasks & Decisions screens still render empty despite the store having the data (Overview reads it) — a per-screen controller/read nuance outstanding; Meetings/GitHub/Approvals/Sync screens not physically populated+tested; Notifications/Search not deep-tested.




screen_overview




screen_tasks




screen_decisions



screen_team



screen_activity



screen_settings

8 GitHub

PARTIAL — connect→read→analyze→plan→approve→modify→PR flow not run (needs explicit GitHub App/creds).

9 Offline / sync

PARTIAL — realtime reconnect confirmed (WS 101); disconnect→queue→conflict→recovery→no-data-loss playbook not adversarially run.


10 Security

PARTIAL — isolation/permission/approval code reviewed in source; live unauthorized-access, private-memory-protection, BYOK-safety, prompt-injection-into-prompt tests not executed.

11 Visual QA (Mimo v2.5)

PASS — flagship chat reviewed: clean baseline logo/profile alignment, no emoji/cartoon decor, dark monochrome + restrained geometry + spectral AI activity; verdict PASS with no required corrections.

PARTIAL — full per-screen Mimo loop (every screen + populated/empty/loading/error state, light/dark) not run yet.



polish_chat

12 Final tech QA

PASS — backend 365 tests, typecheck clean, lint clean; frontend 488 tests, tsc clean; realtime WS 101 no reconnect loop; persistent 404 is the unimplemented /groups/:id/flags endpoint (frontend falls back to DEFAULT_FLAGS).

PARTIAL — production build + full browser console/network sweep for every route not run in this pass.

13 Evidence & verdict

VERIFIED FOUNDATION (real): auth, groups/projects, realtime 'Connected', chat send + @Odin research streamed, all 9 screens render, artifacts+overview populated from real backend, polish verified (Mimo PASS). Backend+frontend suites green.

Honest verdict: the core product + real AI + realtime + real persistence are genuinely working and committed. The deep audit is PARTIAL where explicitly stated: every artifact-viewer geometry, Tasks/Decisions screen-read nuance, Memory loop, GitHub, offline/sync, security adversarial, research auto-trigger, per-screen Mimo loop, and production build are NOT claimed complete.

Known infra finding: the persistent browser session intermittently resets to onboarding on cross-route nav (likely the client-only onboarding-complete flag) — investigate.